

Published research and open source

T4 · Slide 4 and appendix A1 / A4 · Public research evidence

CLAIM	DIRECT SOURCE	SCOPE
RL4CO, KDD 2025; 27 environments; 23 baselines	Publication DOI · Open full paper	Publication and counts supported by the paper and KAIST publication record. The “KDD Oral” designation is not confirmed by the reviewed sources.
Learned-hybrid ~250× faster at 7K nodes	Paper Table 5.4	LKH-3: 501 s, objective 245.0; GLOP-G (AM), RL4CO implementation: 2.0 s, objective 190.9. 501 / 2.0 = 250.5×. Published research comparison, not an Omelet customer-product or cuOpt measurement.
61 ML conference papers; 36 at NeurIPS / ICML / ICLR	KAIST SILAB publication list	Checked 17 September 2026: 61 entries in the ML Conference Papers list; 36 venue lines match the three named conferences. Lab publication history, not a count of company-authored product papers.
897 RL4CO stars in the pitch	Public repository · Dated supplied screenshot	The supplied 14 September screenshot shows 900 stars / 153 forks. Counts change; it does not establish the exact earlier 897 snapshot.

Research versus commercial evidence

The public paper and repository support the research foundation and its authorship. They do not establish the deployment state of Omelet’s commercial platform or NVIDIA components at any named customer. The repository describes an earlier NeurIPS 2023 GLFrontiers Workshop oral; that distinction should not be silently transferred to KDD 2025.

Publication count method

Count direct entries in the lab page’s ML Conference Papers list; then count their venue lines containing NeurIPS, ICML or ICLR. Workshop and engineering-paper lists are excluded. The live list may change after this dated review.

[Open dated publication-count register \(CSV\)](#)